




Parallel Programming Using Compute Unified Device Architecture

ISL Lab Seminar
Han-Sol Kang

INDEX



CU

What is CUDA?



CUDA® is a parallel computing platform and programming model invented by NVIDIA. It enables dramatic increases in computing performance by harnessing the power of the graphics processing unit (GPU).


* http://www.nvidia.com/object/cuda_home_new.html#sthash.1HqbToYO.dpuf

DA


CU

What is CUDA?

CPU Vs. GPU



CPU
Multiple Cores




GPU
Thousand of Cores

DA


CU

What is CUDA?

CPU Vs. GPU



CPU
Multiple Cores




GPU
Thousand of Cores

GPGPU : General Purpose Graphic Processing Unit

DA

CUDA What is CUDA

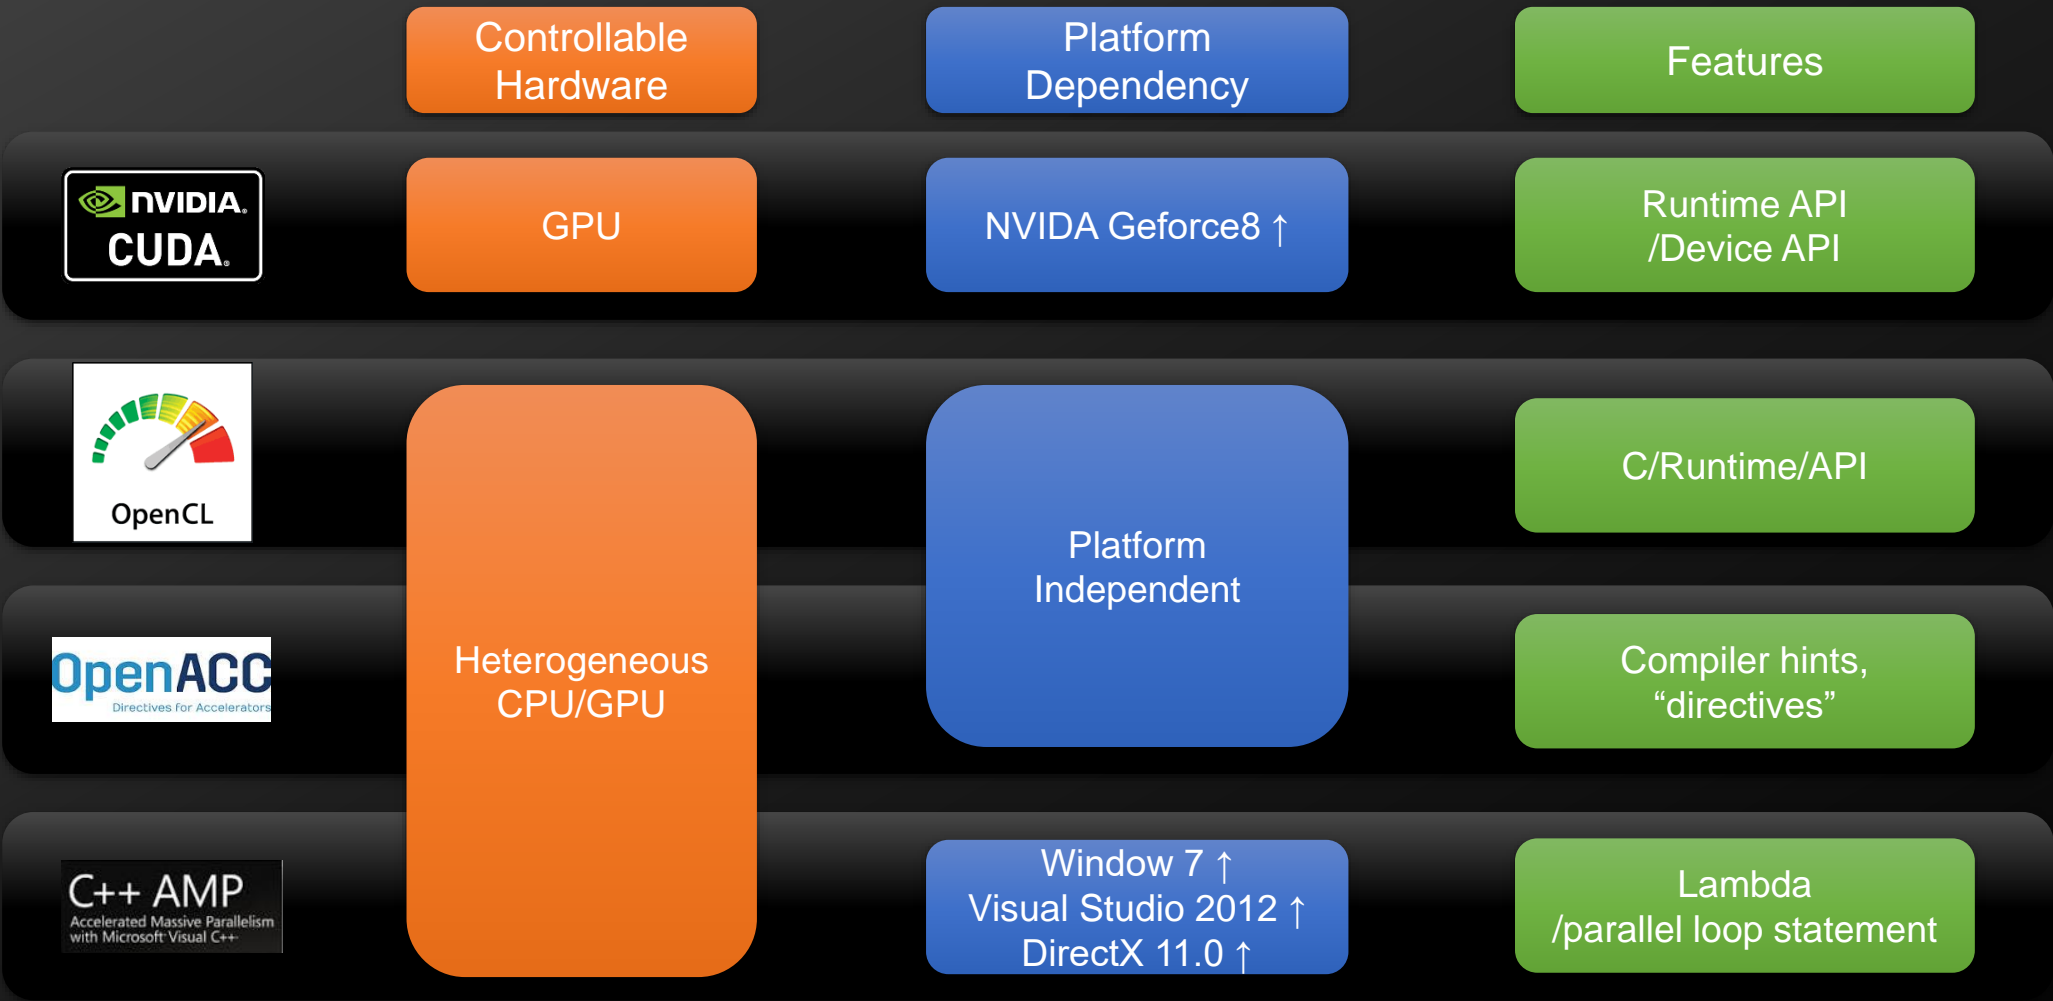
Floating point Operation per second for the CPU and GPU



* <http://docs.nvidia.com/cuda/cuda-c-programming-guide/>

CUDA What is CUDA

GPGPU Programming Models *



* 이현진, et al. "GPGPU 프로그래밍 모델의 기술 동향." 한국정보처리학회 춘계학술대회/ 20.1 (2013): 389-391.

DA

CUDA Installing CUDA 7.0 with VS 2013

ISL Browser

← → ↻

NVIDIA CUDA ZONE [Getting Started](#) [Downloads](#) [Training](#) [Ecosystem](#) [Forums](#) [Register Now](#) [Login](#)

OpenACC Toolkit Now Available

Get started with accelerated directives today

[Learn More>](#)

About CUDA

All about the NVIDIA CUDA parallel computing platform

[Learn more >](#)

Getting Started

First steps for getting started in parallel computing

[Learn more >](#)

Tools & Ecosystem

From accelerated cloud appliances to profiling tools, a gold mine of information

[Learn more >](#)

Academic Collaboration

Partner with NVIDIA to advance parallel computing education and research

[Learn more >](#)

CUDA Downloads

Get the latest and greatest version of the CUDA Toolkit

[Learn more >](#)

Resources

Materials and links especially for GPU Computing professionals and developers

[Learn more >](#)

CUDA Week In Review

[Subscribe me](#)

GPU Computing

[Follow](#)

Alex Z. @alexzitti 30 Aug
@GPUComputing The power of cuDNN is all around us, stay tuned ;) #nvidia #cuda #cudnn #neuralnetwork
Retweeted by CUDA, GPU Computing
[Expand](#)

Mike Wang @mikepcw 27 Aug
Prof Paul Bonnington of @monashuni introduces the microscope of 21st century #gputechconf #singapore

Click

DA

8

ISL Image System Laboratory

CUDA Installing CUDA 7.0 with VS 2013

ISL Browser

← → ↻

NVIDIA CUDA ZONE [Getting Started](#) [Downloads](#) [Training](#) [Ecosystem](#) [Forums](#) [Register Now](#) [Login](#)

[Home](#) > [CUDA ZONE](#) > [Tools & Ecosystem](#) > [CUDA Toolkit](#) > [CUDA 7 Downloads](#)

CUDA 7 Downloads

Check out:

- [CUDA 7 Performance Report and Webinar Recording](#)
- [An informative webinar by Ujval Kapasi, NVIDIA's CUDA Product Manager CUDA 7 Features and Overview](#)
- [The Power of C++11 in CUDA 7](#), another technical blog on Parallel Forall.

If you find any issues please [file a bug](#) (requires membership of the [CUDA Registered Developer Program](#)).

Please Note: There is a recommended patch for CUDA 7.0 which resolves an issue in the cuFFT library that can lead to incorrect results for certain inputs sizes less than or equal to 1920 in any dimension when cuFFTSetStream() is passed a [non-blocking stream](#) (e.g., one created using the cudaStreamNonBlocking flag of the CUDA Runtime API or the CU_STREAM_NON_BLOCKING flag of the CUDA Driver API).

CUDA Toolkit 7.5 Release Candidate Now Available For All Developers [Learn More](#).

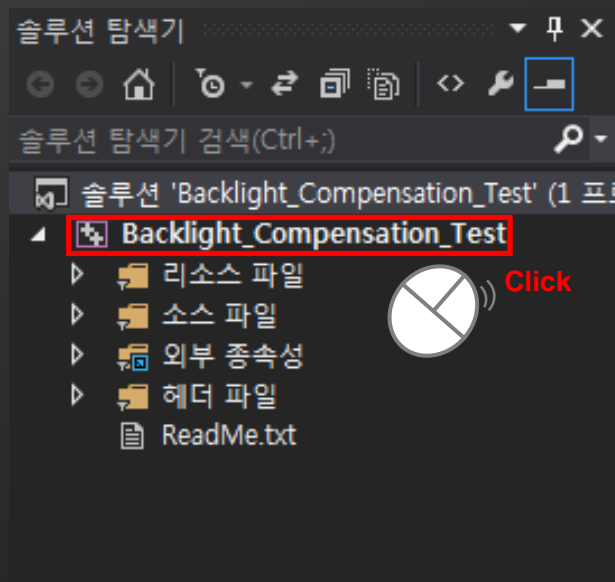
WindowsLinux x86Linux POWER8Mac OSX

Version	Network Installer	Click Local Installer
Windows 8.1	EXE (8.0MB)	EXE (939MB)
Windows 7		
Win Server 2012 R2		
Win Server 2008 R2		

Documentation
Release Notes
End User License Agreement
Online Documentation
CUDA Toolkit Overview

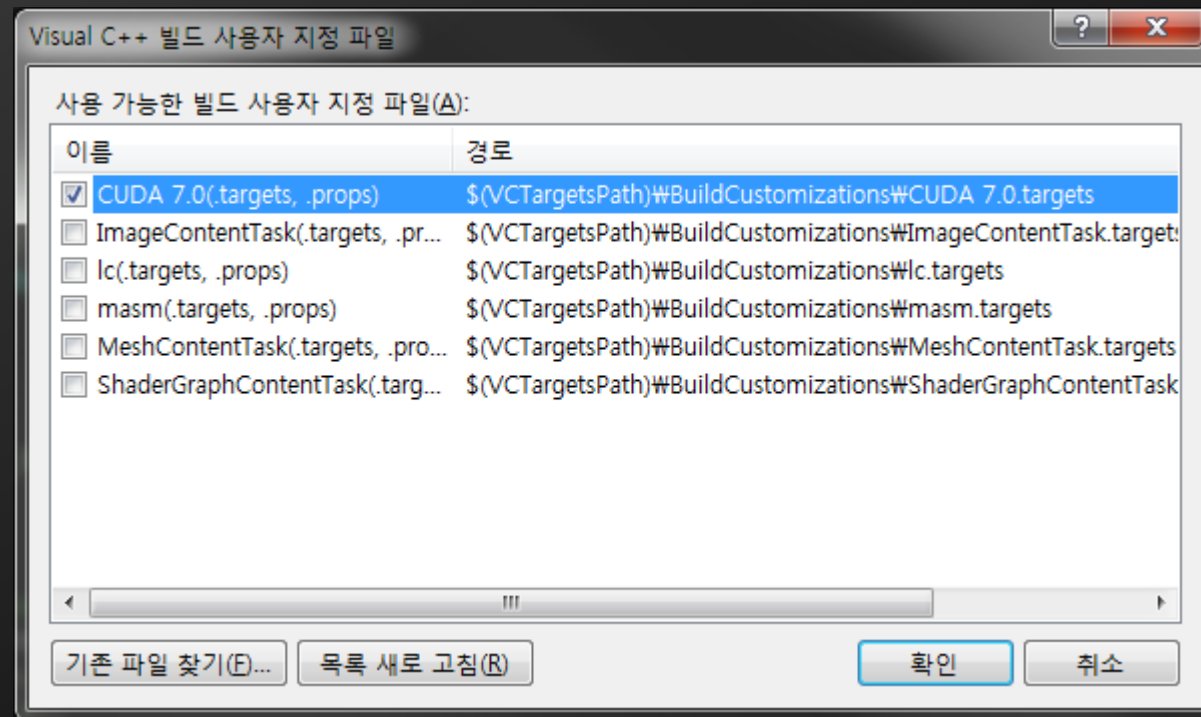
CUDA Installing CUDA 7.0 with VS 2013

Visual Studio Settings



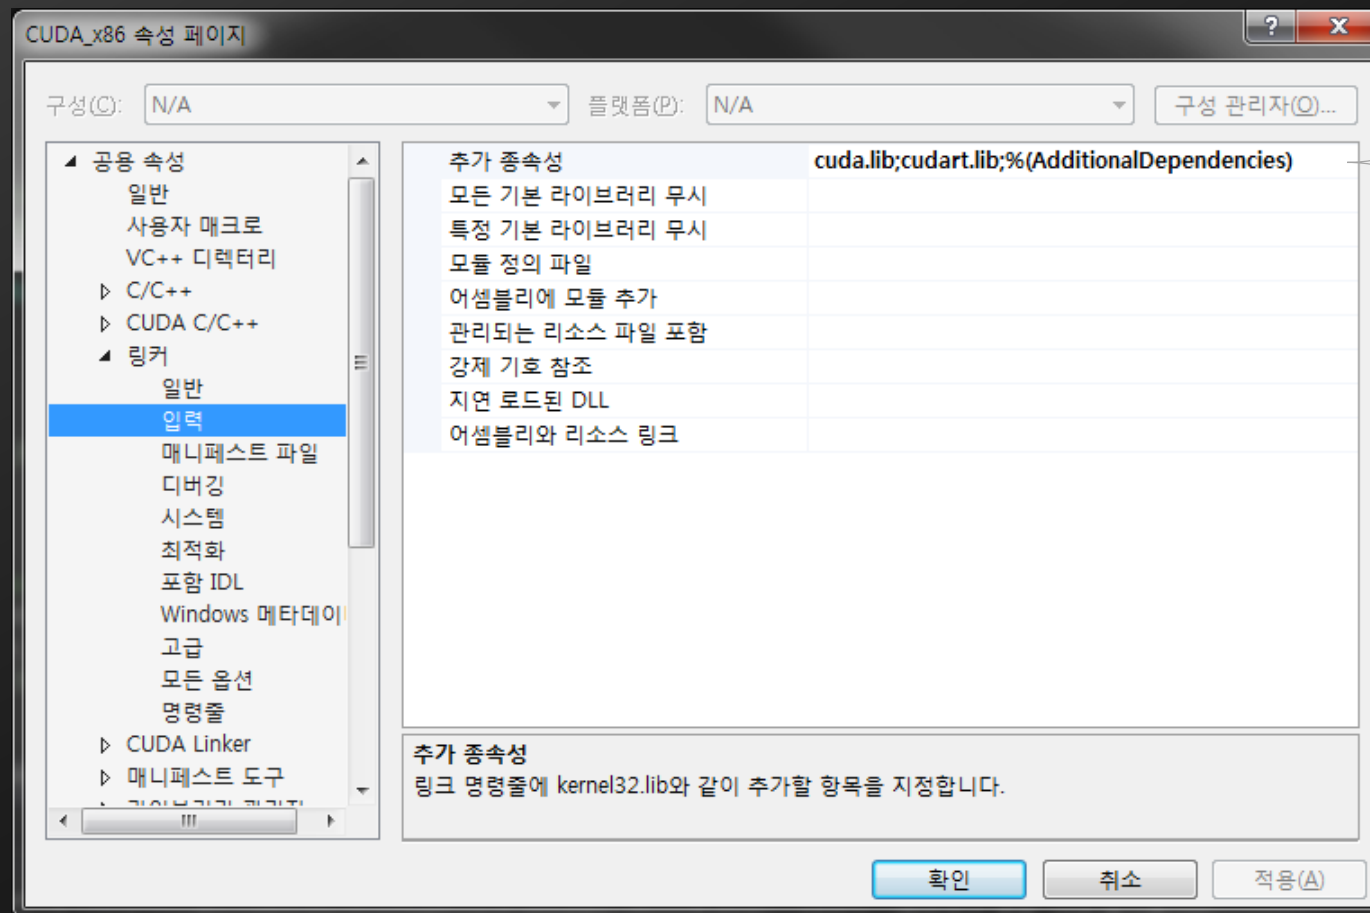
CUDA Installing CUDA 7.0 with VS 2013

Visual Studio Settings



CUDA Installing CUDA 7.0 with VS 2013

Visual Studio Settings



cublas.lib
 cublas_device.lib
 cuda.lib
 cudadevrt.lib
 cudart.lib
 cudart_static.lib
 cufft.lib
 cufftw.lib
 curand.lib
 cusparse.lib
 nppc.lib
 nppi.lib
 npps.lib
 nvblas.lib
 nvcudvid.lib
 nvrtc.lib
 OpenCL.lib

CUDA Installing CUDA 7.0 with VS 2013

Unsupported Features


- ◆ 32-bit x86 Linux Systems
- ◆ Red Hat Enterprise Linux 5 and CentOS5
- ◆ Tesla Architecture
- ◆ Certain CUDA Features on 32-bit and 32-on-64-bit Windows Systems

Running 32-bit applications on Tesla and Quadro products

Using the Thrust library form 32-bit applications

32-bit version of the CUDA Toolkit scientific libraries, including cuBILAS, cuSPARSE, cuFFT, cuRAND and NPP


32-bit version of the CUDA samples



CUBasic of CUDA

Terminology

Host : The CPU and its memory



Device : The GPU and its memory



Fermi Architecture



CUDA Basic of CUDA

Directives

<code>__global__</code>	Host to Device Return type is void, Not allow Recursion
<code>__device__</code>	Device to Device
<code>__host__</code>	Host to Host

Graphic Card Memory Allocation


```
cudaError_t cudaMalloc(void** devPtr, size_t cout)
```

```
cudaError_t cudaFree(void* devPtr)
```

```
cudaError_t cudaMemcpy(void* dst, const void* src, size_t cout, cudaMemcpyKind kind)
```

CU Basic of CUDA

Thread, Block, Grid



CU

Basic of CUDA

Kernel Function

```
__global__ void kernel << <Dg, Db, Ns, S >> > ();
```

Dg(Dimensions of the grid) : type dim3


Db(Dimensions of the block) : type dim3

Ns(Number of bytes shared memory dynamically allocated /block) : type size_t

S(associated cudaStream_t)

ex) `__global__ void kernel << <1, 512 >> > (int a, int b, int c);`

G80



CU

Basic of CUDA

Kernel Function

```
__global__ void kernel << <Dg, Db, Ns, S >> > ();
```

Dg(Dimensions of the grid) : type dim3

Db(Dimensions of the block) : type dim3

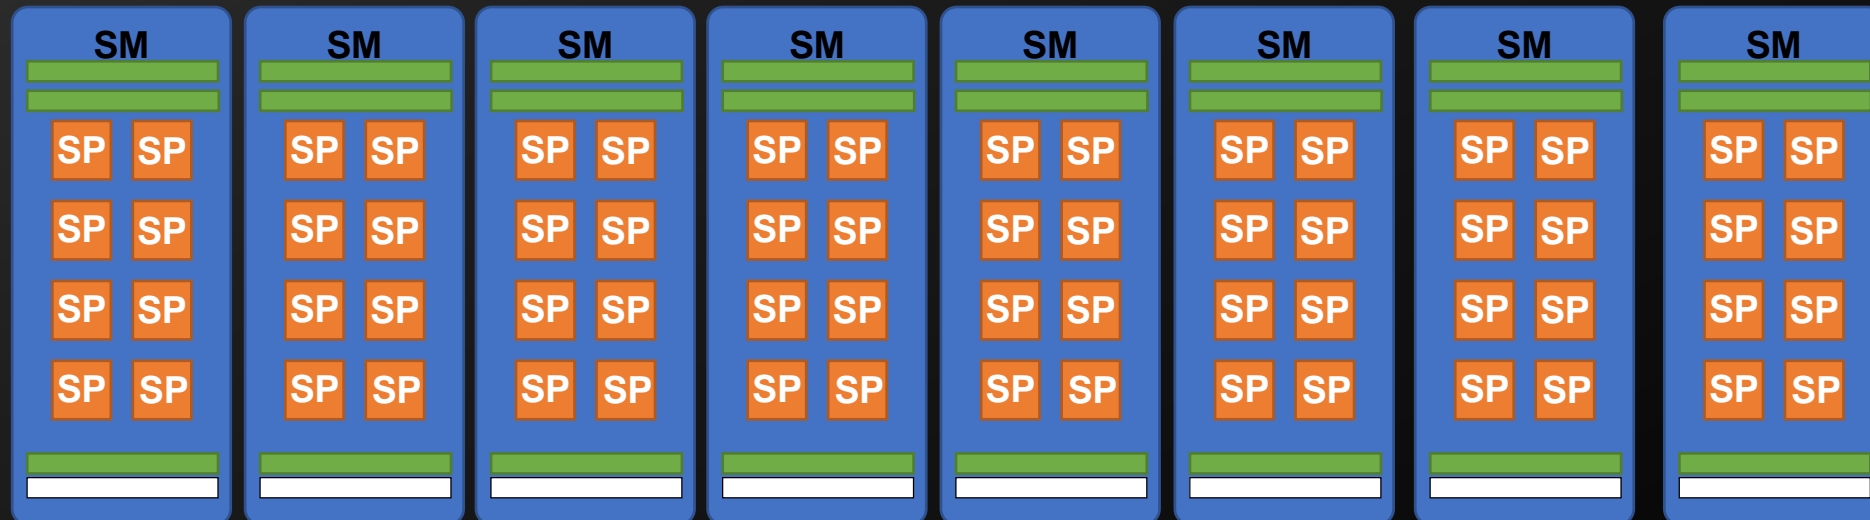
Ns(Number of bytes shared memory dynamically allocated /block) : type size_t

S(associated cudaStream_t)

ex)

```
__global__ void kernel << <8, 64 >> > (int a, int, b, int c);
```

G80



Basic of CUDA

Example

```
__global__ void VectorAdd(int *a, int *b, int *c, int n)
{
    int i = threadIdx.x;
    if (i < n)
    {
        c[i] = a[i] + b[i];
    }
}
```

```
int main()
{
    int *a, *b, *c;
    int *d_a, *d_b, *d_c;

    a = (int *)malloc(SIZE*sizeof(int));
    b = (int *)malloc(SIZE*sizeof(int));
    c = (int *)malloc(SIZE*sizeof(int));

    cudaMalloc(&d_a, SIZE*sizeof(int));
    cudaMalloc(&d_b, SIZE*sizeof(int));
    cudaMalloc(&d_c, SIZE*sizeof(int));
    for (int i = 0; i < SIZE; ++i)
    {
        a[i] = i;
        b[i] = i;
        c[i] = 0;
    }
}
```

```
cudaMemcpy(d_a, a, SIZE*sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, SIZE*sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(d_c, c, SIZE*sizeof(int), cudaMemcpyHostToDevice);
```

```
VectorAdd <<< 1, SIZE >>>(d_a, d_b, d_c, SIZE);
cudaMemcpy(c, d_c, SIZE*sizeof(int), cudaMemcpyDeviceToHost);
for (int i = 0; i < 10; ++i)
{
    printf("c[%d] = %d\n", i, c[i]);
}
```

```
free(a);
free(b);
free(c);
```

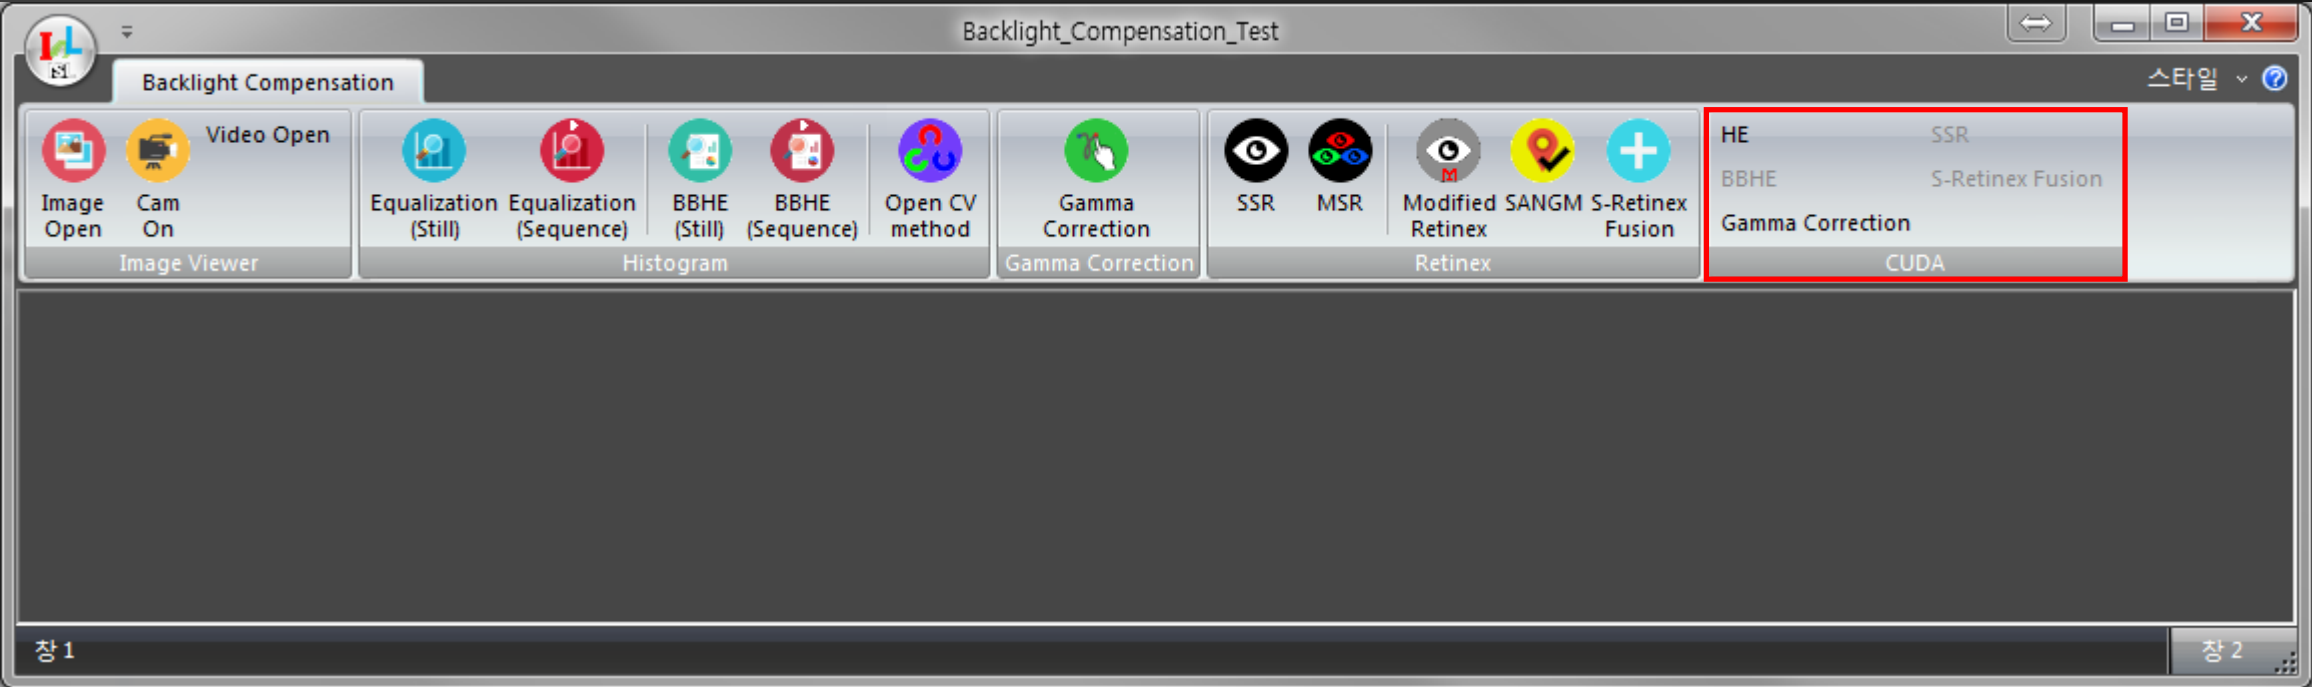
```
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
return 0;
```

```
c[0] = 0
c[1] = 2
c[2] = 4
c[3] = 6
c[4] = 8
c[5] = 10
c[6] = 12
c[7] = 14
c[8] = 16
c[9] = 18
```

계속하려면 아무 키나 누르십시오 . . .

CU Applications

Backlight Compensation Test Program



DA

Applications

Classes of BCT

CHistManager

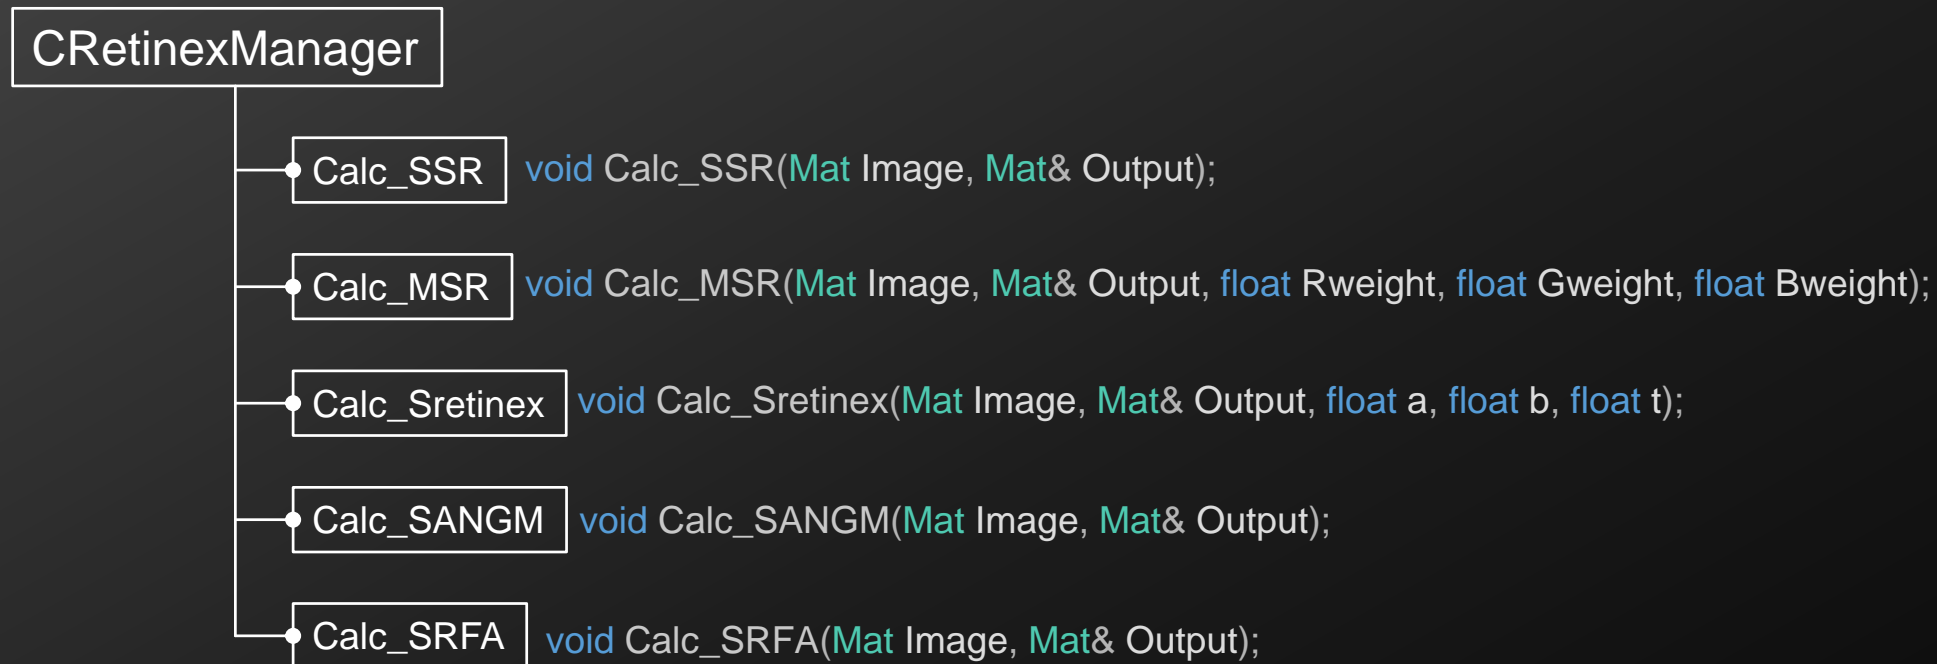
- Calc_Histogram `void Calc_Histogram(Mat Image, float* fHist, int colorcode);`
- Calc_Hist_Eq `void Calc_Hist_Eq(Mat Image, Mat& Output, float* fHist, float* Trans, int colorcode, int nStart, int nEnd);`
- Calc_Hist_Eq3 `void Calc_Hist_Eq3(Mat Image, Mat& Output, float* fHist, float* Trans, int colorcode, int nStart, int nEnd);`
- Hist_Visual `void Hist_Visual(Mat Image, Mat& Output, int colorcode);`

CAvgStdGamma

- Calc_Avg_Std `void Calc_Avg_Std(Mat Image, float &average, float &sigma);`
- Gamma_Correction `void Gamma_Correction(Mat Image, Mat& Output, float gamma);`

Applications

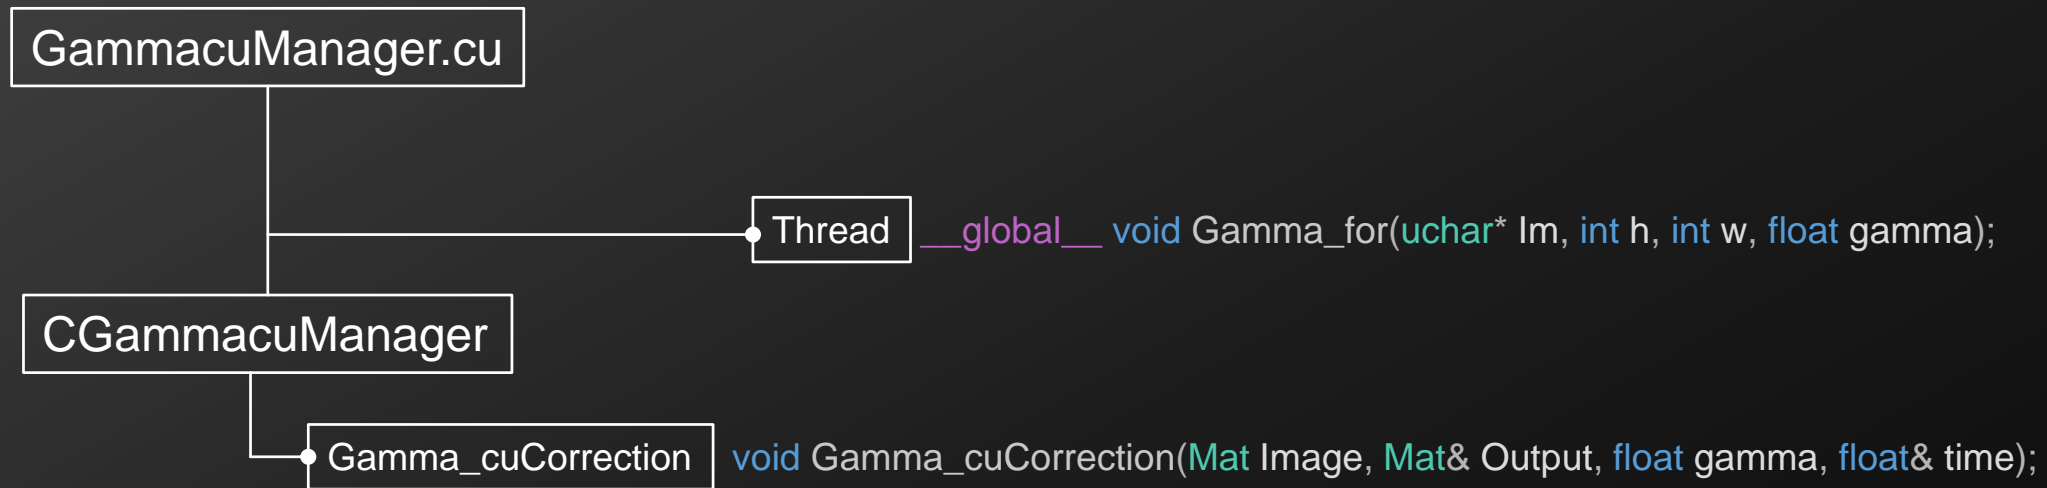
Classes of BCT



CU

Applications

Classes of BCT

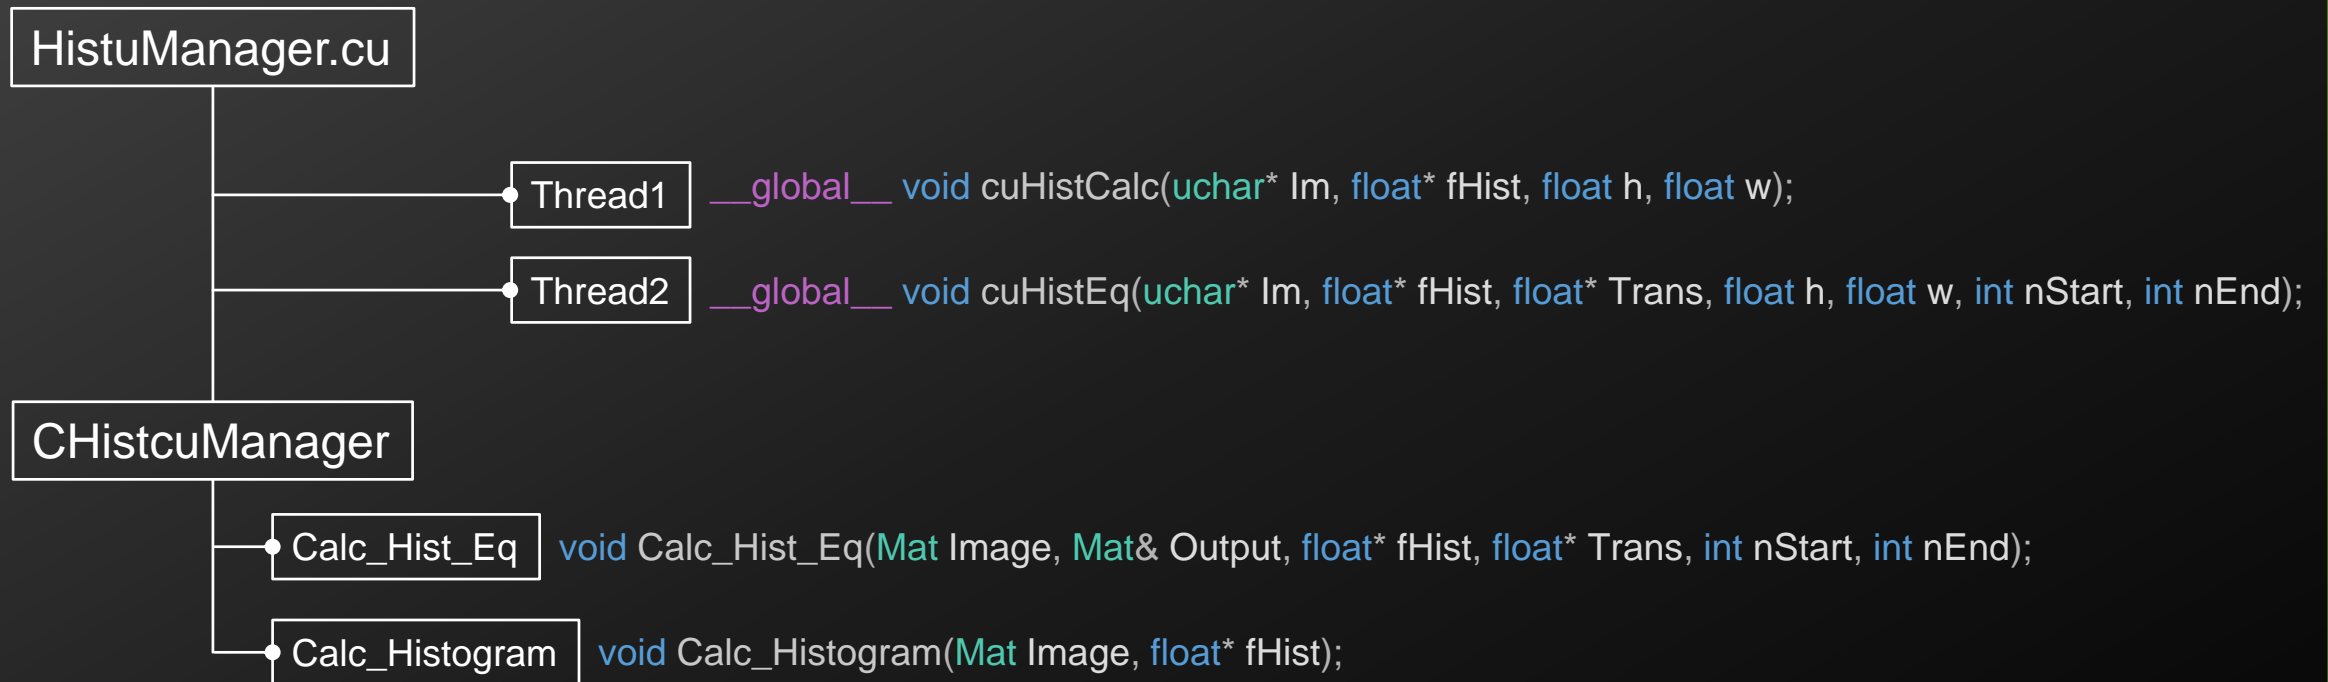


DA

CU

Applications

Classes of BCT



CU

Applications

Example – Histogram

```

void CHistcuManager::Calc_Histogram(Mat Image, float* fHist)
{
    float h = Image.size().height;
    float w = Image.size().width;
    uchar *Im = Image.ptr<uchar>(0);

    uchar *d_Im;
    float *d_fHist;
    size_t SIZE = h*w;

    cudaMalloc((void**)&d_Im, SIZE*sizeof(uchar));
    cudaMalloc((void**)&d_fHist, (size_t)256 * sizeof(float));

    cudaMemcpy(d_Im, Im, SIZE*sizeof(uchar), cudaMemcpyHostToDevice);
    cudaMemcpy(d_fHist, fHist, (size_t)256 * sizeof(float), cudaMemcpyHostToDevice);
    cuHistCalc <<< 65535, 512 >>> (d_Im, d_fHist, h, w);

    cudaMemcpy(Im, d_Im, SIZE*sizeof(uchar), cudaMemcpyDeviceToHost);
    cudaMemcpy(fHist, d_fHist, (size_t)256 * sizeof(float), cudaMemcpyDeviceToHost);
    cudaFree(d_Im);
    cudaFree(d_fHist);
}

```

```

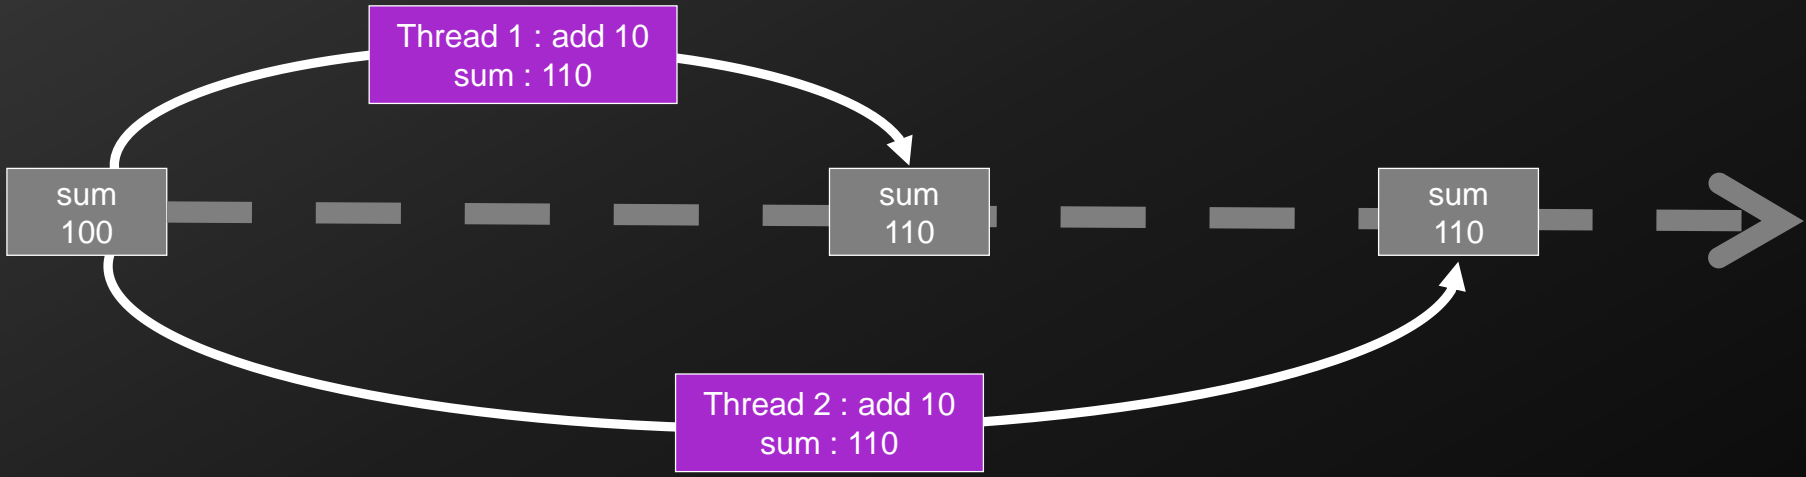
__global__ void cuHistCalc(uchar* Im, float* fHist, float h, float w)
{
    int i = blockDim.x*blockIdx.x + threadIdx.x;

    if (i < h*w)
    {
        atomicAdd(&(fHist[Im[i]]), 1);
    }
}

```

CU Applications


Race Condition



DA

CU Applications

Race Condition



DA

CU

Applications

Atomic Operation

Arithmetic

`atomicAdd()` computes $(old + val)$, and stores the result back to memory at the same address

`atomicSub()` computes $(old - val)$, and stores the result back to memory at the same address

`atomicExch()` computes $swap(old, val)$, and stores the result back to memory at the same address

`atomicMin()` computes $Min(old, val)$, and stores the result back to memory at the same address

`atomicMax()` computes $Max(old, val)$, and stores the result back to memory at the same address

`atomicInc()` computes $((old \geq val) ? val : (old-1))$, and stores the result back to memory at the same address

`atomicDec()` computes $((old == 0) \vee (old > val) ? val : (old-1))$, and stores the result back to memory at the same address

`atomicCAS()` computes $(old == compare ? val : old)$, and stores the result back to memory at the same address

CU

Applications

Atomic Operation

Bit

`atomicAnd()` (old & val), and stores the result back to memory at the same address

`atomicOr()` (old | val), and stores the result back to memory at the same address

`atomicXor()` (old ^ val), and stores the result back to memory at the same address

DA

CU

Applications

Example – Gamma Correction

```

void CGammacuManager::Gamma_cuCorrection(Mat Image, Mat& Output, float gamma)
{
    float elapsed_time_ms = 0.0f;
    cudaEvent_t start, stop;

    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);

    cvtColor(Image, Image, CV_RGB2YCrCb);           //YCrCb
    split(Image, m_vecChannel);

    int h = Image.size().height;
    int w = Image.size().width;
    uchar* Im = m_vecChannel[0].ptr<uchar>(0);
    uchar* dev_Im;
    cudaMalloc(&dev_Im, h*w*sizeof(uchar));
    cudaMemcpy(dev_Im, Im, h*w*sizeof(uchar), cudaMemcpyHostToDevice);

    Gamma_for <<<65535, 512>>> (dev_Im, h, w, gamma);

    cudaMemcpy(Im, dev_Im, h*w*sizeof(uchar), cudaMemcpyDeviceToHost);
    merge(m_vecChannel, Output);
    cvtColor(Output, Output, CV_YCrCb2RGB);
    cudaFree(dev_Im);
}

```

```

__global__ void Gamma_for(uchar* Im, int h, int w, float gamma)
{
    const int idx = blockIdx.x * blockDim.x + threadIdx.x;

    float inv_gamma = 1.f / gamma;
    if ( (idx < w*h) )
    {
        Im[idx] = Clipping(pow((Im[idx] / 255.f), inv_gamma) * 255.f + 0.5f);
    }
}

```

CU Applications

Gamma Correction Results

기준영상
(1920 x 1080)



Entire (unit : s)		
Trial	Normal Function	CUDA Function
1	0.102687	0.014429
2	0.101979	0.015645
3	0.102959	0.015179
4	0.109468	0.014846
5	0.104883	0.014513
6	0.102613	0.014487
7	0.101135	0.01484
8	0.101176	0.014149
9	0.104109	0.014452
10	0.101946	0.01418
11	0.102001	0.015072
12	0.101744	0.015184
13	0.102265	0.0143
14	0.100986	0.014198
15	0.101596	0.014725
16	0.101665	0.014718
17	0.102031	0.014338
18	0.101583	0.01496
19	0.102474	0.014493
20	0.101204	0.0014437
21	0.100953	0.014219
22	0.101229	0.014535
23	0.101547	0.014537
24	0.10177	0.014708
25	0.10201	0.014421
26	0.101805	0.014247
27	0.101574	0.01591
28	0.101128	0.014228
29	0.103729	0.014163
30	0.100847	0.01403
Average 0.102236533 0.014171657		



Q

&

A

Thank you for your attention